

# Load Data and Import Dependencies

In [2]: `!pip install mglearn`

```
Requirement already satisfied: mglearn in c:\users\administrator\anaconda3\lib\site-packages (0.2.0)
Requirement already satisfied: numpy in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (1.26.4)
Requirement already satisfied: matplotlib in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (3.9.2)
Requirement already satisfied: scikit-learn in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (1.5.1)
Requirement already satisfied: pandas in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (2.2.2)
Requirement already satisfied: pillow in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (10.4.0)
Requirement already satisfied: cycler in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (0.11.0)
Requirement already satisfied: imageio in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (2.33.1)
Requirement already satisfied: joblib in c:\users\administrator\anaconda3\lib\site-packages (from mglearn) (1.4.2)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\administrator\anaconda3\lib\site-packages (from matplotlib->mglearn) (1.2.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\administrator\anaconda3\lib\site-packages (from matplotlib->mglearn) (4.51.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\administrator\anaconda3\lib\site-packages (from matplotlib->mglearn) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\administrator\anaconda3\lib\site-packages (from matplotlib->mglearn) (24.1)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\administrator\anaconda3\lib\site-packages (from matplotlib->mglearn) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\administrator\anaconda3\lib\site-packages (from matplotlib->mglearn) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\administrator\anaconda3\lib\site-packages (from pandas->mglearn) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\administrator\anaconda3\lib\site-packages (from pandas->mglearn) (2023.3)
Requirement already satisfied: scipy>=1.6.0 in c:\users\administrator\anaconda3\lib\site-packages (from scikit-learn->mglearn) (1.13.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\administrator\anaconda3\lib\site-packages (from scikit-learn->mglearn) (3.5.0)
Requirement already satisfied: six>=1.5 in c:\users\administrator\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib->mglearn) (1.16.0)
```

In [3]: `import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split`

```
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
```

```
In [4]: df = pd.read_csv('data/KaggleV2-May-2016.csv')
df
```

```
Out[4]:
```

|        | PatientId    | AppointmentID | Gender | ScheduledDay         | AppointmentDay       | Age | Ne  |
|--------|--------------|---------------|--------|----------------------|----------------------|-----|-----|
| 0      | 2.987250e+13 | 5642903       | F      | 2016-04-29T18:38:08Z | 2016-04-29T00:00:00Z | 62  |     |
| 1      | 5.589978e+14 | 5642503       | M      | 2016-04-29T16:08:27Z | 2016-04-29T00:00:00Z | 56  |     |
| 2      | 4.262962e+12 | 5642549       | F      | 2016-04-29T16:19:04Z | 2016-04-29T00:00:00Z | 62  | M   |
| 3      | 8.679512e+11 | 5642828       | F      | 2016-04-29T17:29:31Z | 2016-04-29T00:00:00Z | 8   |     |
| 4      | 8.841186e+12 | 5642494       | F      | 2016-04-29T16:07:23Z | 2016-04-29T00:00:00Z | 56  |     |
| ...    | ...          | ...           | ...    | ...                  | ...                  | ... | ... |
| 110522 | 2.572134e+12 | 5651768       | F      | 2016-05-03T09:15:35Z | 2016-06-07T00:00:00Z | 56  |     |
| 110523 | 3.596266e+12 | 5650093       | F      | 2016-05-03T07:27:33Z | 2016-06-07T00:00:00Z | 51  |     |
| 110524 | 1.557663e+13 | 5630692       | F      | 2016-04-27T16:03:52Z | 2016-06-07T00:00:00Z | 21  |     |
| 110525 | 9.213493e+13 | 5630323       | F      | 2016-04-27T15:09:23Z | 2016-06-07T00:00:00Z | 38  |     |
| 110526 | 3.775115e+14 | 5629448       | F      | 2016-04-27T13:30:56Z | 2016-06-07T00:00:00Z | 54  |     |

110527 rows × 14 columns

## Analysis and Visualization

```
In [6]: df.shape
```

```
Out[6]: (110527, 14)
```

```
In [7]: df.describe()
```

Out[7]:

|              | PatientId    | AppointmentID | Age           | Scholarship   | Hipertension  | D      |
|--------------|--------------|---------------|---------------|---------------|---------------|--------|
| <b>count</b> | 1.105270e+05 | 1.105270e+05  | 110527.000000 | 110527.000000 | 110527.000000 | 110527 |
| <b>mean</b>  | 1.474963e+14 | 5.675305e+06  | 37.088874     | 0.098266      | 0.197246      | 0      |
| <b>std</b>   | 2.560949e+14 | 7.129575e+04  | 23.110205     | 0.297675      | 0.397921      | 0      |
| <b>min</b>   | 3.921784e+04 | 5.030230e+06  | -1.000000     | 0.000000      | 0.000000      | 0      |
| <b>25%</b>   | 4.172614e+12 | 5.640286e+06  | 18.000000     | 0.000000      | 0.000000      | 0      |
| <b>50%</b>   | 3.173184e+13 | 5.680573e+06  | 37.000000     | 0.000000      | 0.000000      | 0      |
| <b>75%</b>   | 9.439172e+13 | 5.725524e+06  | 55.000000     | 0.000000      | 0.000000      | 0      |
| <b>max</b>   | 9.999816e+14 | 5.790484e+06  | 115.000000    | 1.000000      | 1.000000      | 1      |

In [8]: `df.head`

```
Out[8]: <bound method NDFrame.head of
```

|        | ScheduledDay \ |         | PatientId | AppointmentID        | Gender |
|--------|----------------|---------|-----------|----------------------|--------|
| 0      | 2.987250e+13   | 5642903 | F         | 2016-04-29T18:38:08Z |        |
| 1      | 5.589978e+14   | 5642503 | M         | 2016-04-29T16:08:27Z |        |
| 2      | 4.262962e+12   | 5642549 | F         | 2016-04-29T16:19:04Z |        |
| 3      | 8.679512e+11   | 5642828 | F         | 2016-04-29T17:29:31Z |        |
| 4      | 8.841186e+12   | 5642494 | F         | 2016-04-29T16:07:23Z |        |
| ...    | ...            | ...     | ...       | ...                  | ...    |
| 110522 | 2.572134e+12   | 5651768 | F         | 2016-05-03T09:15:35Z |        |
| 110523 | 3.596266e+12   | 5650093 | F         | 2016-05-03T07:27:33Z |        |
| 110524 | 1.557663e+13   | 5630692 | F         | 2016-04-27T16:03:52Z |        |
| 110525 | 9.213493e+13   | 5630323 | F         | 2016-04-27T15:09:23Z |        |
| 110526 | 3.775115e+14   | 5629448 | F         | 2016-04-27T13:30:56Z |        |

|        | AppointmentDay       | Age | Neighbourhood     | Scholarship \ |
|--------|----------------------|-----|-------------------|---------------|
| 0      | 2016-04-29T00:00:00Z | 62  | JARDIM DA PENHA   | 0             |
| 1      | 2016-04-29T00:00:00Z | 56  | JARDIM DA PENHA   | 0             |
| 2      | 2016-04-29T00:00:00Z | 62  | MATA DA PRAIA     | 0             |
| 3      | 2016-04-29T00:00:00Z | 8   | PONTAL DE CAMBURI | 0             |
| 4      | 2016-04-29T00:00:00Z | 56  | JARDIM DA PENHA   | 0             |
| ...    | ...                  | ... | ...               | ...           |
| 110522 | 2016-06-07T00:00:00Z | 56  | MARIA ORTIZ       | 0             |
| 110523 | 2016-06-07T00:00:00Z | 51  | MARIA ORTIZ       | 0             |
| 110524 | 2016-06-07T00:00:00Z | 21  | MARIA ORTIZ       | 0             |
| 110525 | 2016-06-07T00:00:00Z | 38  | MARIA ORTIZ       | 0             |
| 110526 | 2016-06-07T00:00:00Z | 54  | MARIA ORTIZ       | 0             |

|        | Hipertension | Diabetes | Alcoholism | Handcap | SMS_received | No-show |
|--------|--------------|----------|------------|---------|--------------|---------|
| 0      | 1            | 0        | 0          | 0       | 0            | No      |
| 1      | 0            | 0        | 0          | 0       | 0            | No      |
| 2      | 0            | 0        | 0          | 0       | 0            | No      |
| 3      | 0            | 0        | 0          | 0       | 0            | No      |
| 4      | 1            | 1        | 0          | 0       | 0            | No      |
| ...    | ...          | ...      | ...        | ...     | ...          | ...     |
| 110522 | 0            | 0        | 0          | 0       | 1            | No      |
| 110523 | 0            | 0        | 0          | 0       | 1            | No      |
| 110524 | 0            | 0        | 0          | 0       | 1            | No      |
| 110525 | 0            | 0        | 0          | 0       | 1            | No      |
| 110526 | 0            | 0        | 0          | 0       | 1            | No      |

[110527 rows x 14 columns]>

```
In [9]: print(df.isnull().sum())
```

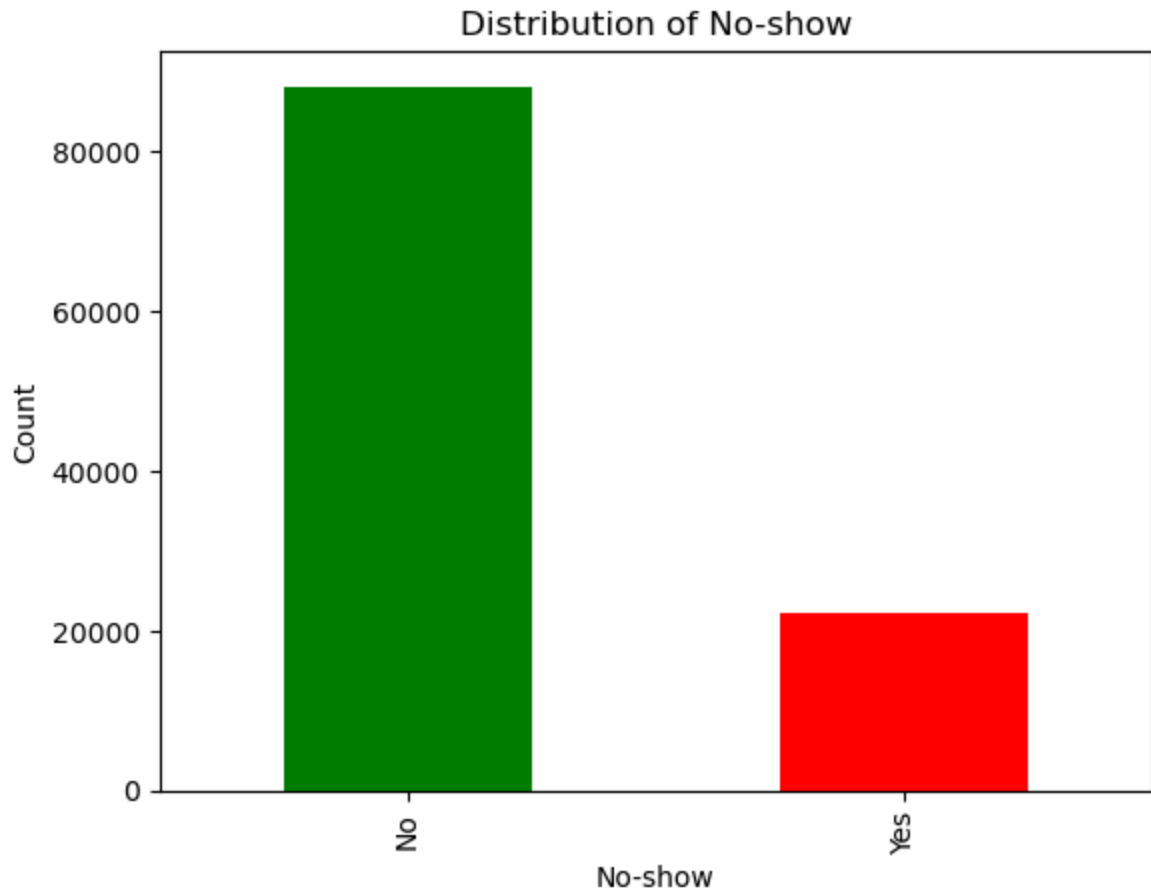
```
PatientId      0
AppointmentID  0
Gender          0
ScheduledDay   0
AppointmentDay  0
Age            0
Neighbourhood  0
Scholarship    0
Hipertension   0
Diabetes        0
Alcoholism     0
Handcap        0
SMS_received   0
No-show        0
dtype: int64
```

```
In [10]: df.dtypes
```


```
Out[10]: PatientId      float64
AppointmentID    int64
Gender           object
ScheduledDay     object
AppointmentDay    object
Age             int64
Neighbourhood    object
Scholarship      int64
Hipertension     int64
Diabetes         int64
Alcoholism       int64
Handcap          int64
SMS_received     int64
No-show          object
dtype: object
```

```
In [11]: # df['Gender'] = pd.to_numeric(df['Gender'], downcast='integer', errors='coerce')
# df.dtypes
```

```
In [12]: # Plotting the distribution of No-show
df['No-show'].value_counts().plot(kind='bar', color=['green', 'red'])
plt.title('Distribution of No-show')
plt.xlabel('No-show')
plt.ylabel('Count')
plt.show()
```

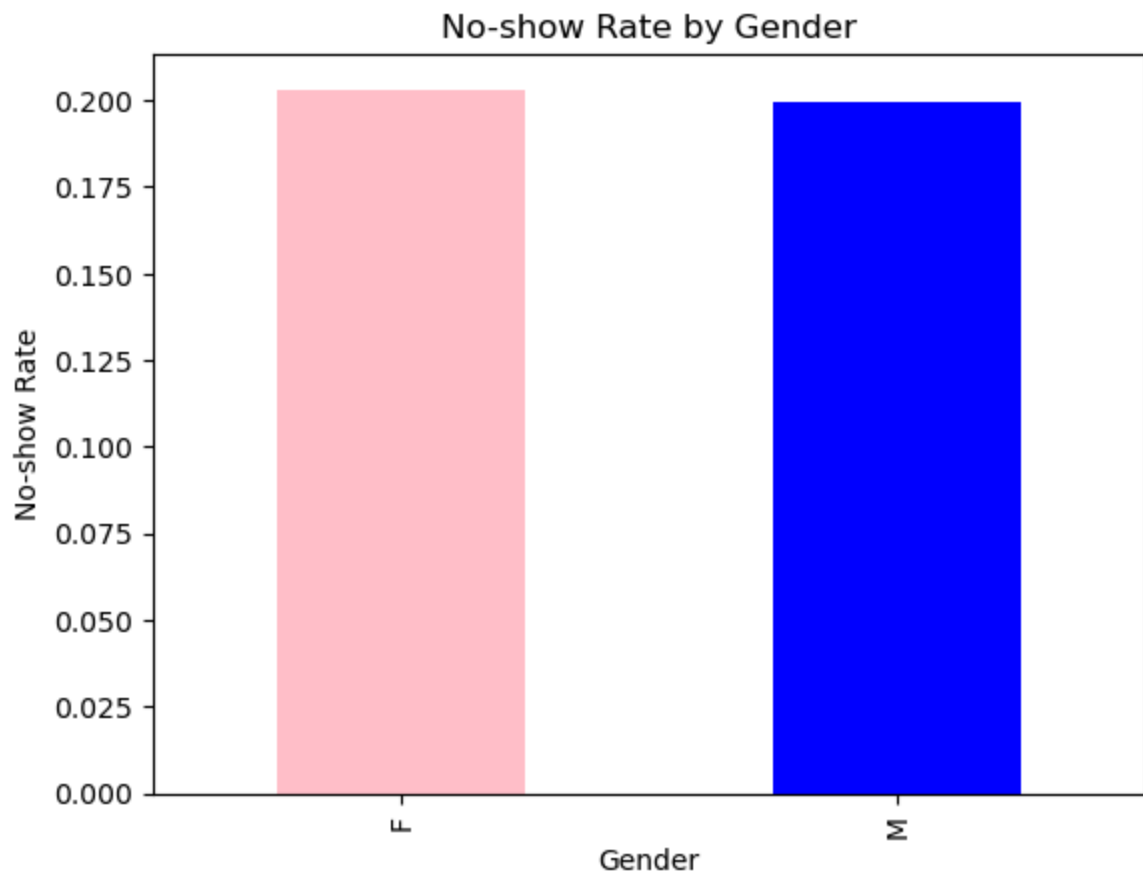


```
In [13]: # Plotting the age distribution
plt.hist(df['Age'], bins=20, color='blue', edgecolor='black')
plt.title('Age Distribution of Patients')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



```
In [14]: # Calculating no-show rate by gender
no_show_rate = df.groupby('Gender')['No-show'].apply(lambda x: (x == 'Yes').mean())

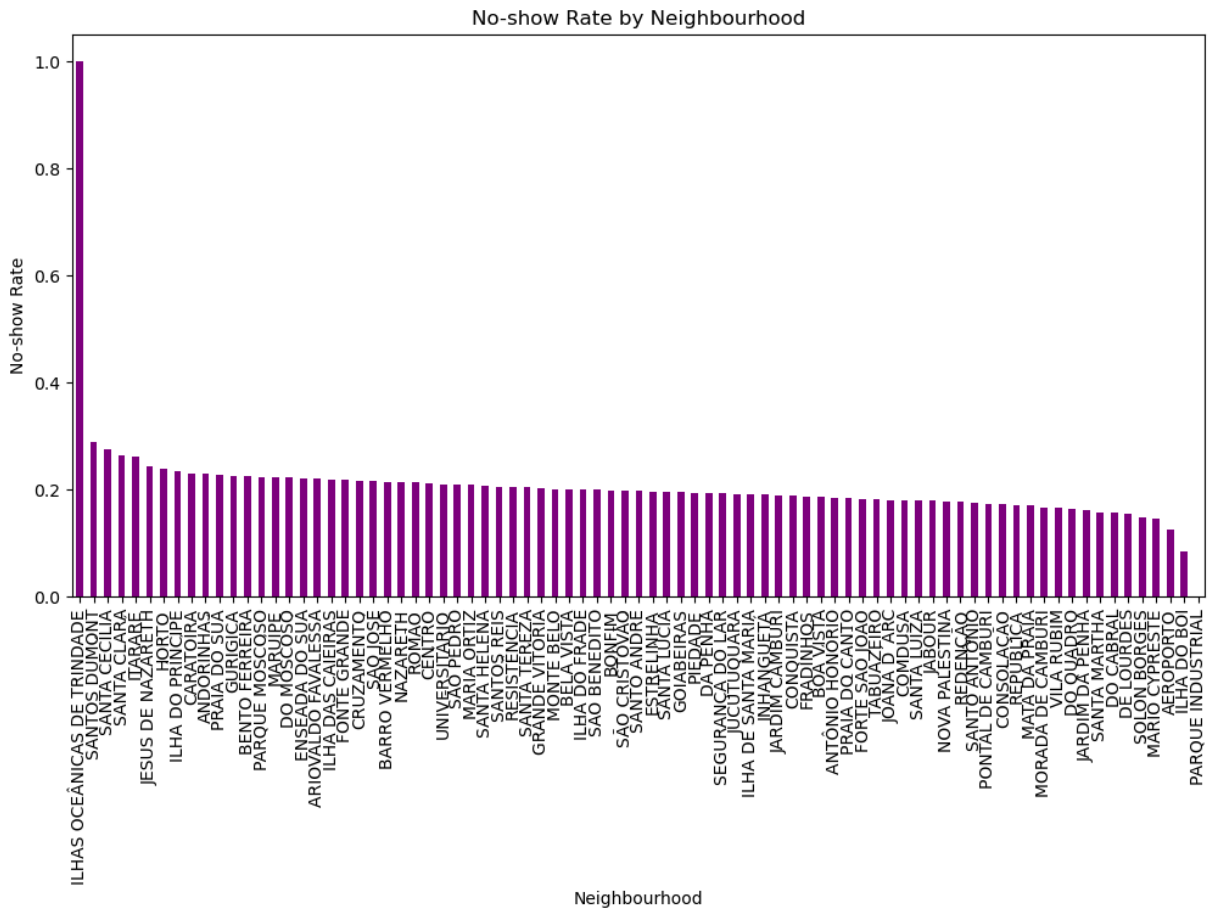
# Plotting the no-show rate by gender
no_show_rate.plot(kind='bar', color=['pink', 'blue'])
plt.title('No-show Rate by Gender')
plt.xlabel('Gender')
plt.ylabel('No-show Rate')
plt.show()
```



```
In [15]: # Calculating no-show rate by neighbourhood
no_show_rate_neighbourhood = df.groupby('Neighbourhood')['No-show'].apply(lambda x:


# Plotting the no-show rate by neighbourhood
no_show_rate_neighbourhood.sort_values(ascending=False).plot(kind='bar', figsize=(1
plt.title('No-show Rate by Neighbourhood')
plt.xlabel('Neighbourhood')
plt.ylabel('No-show Rate')
plt.show()
```





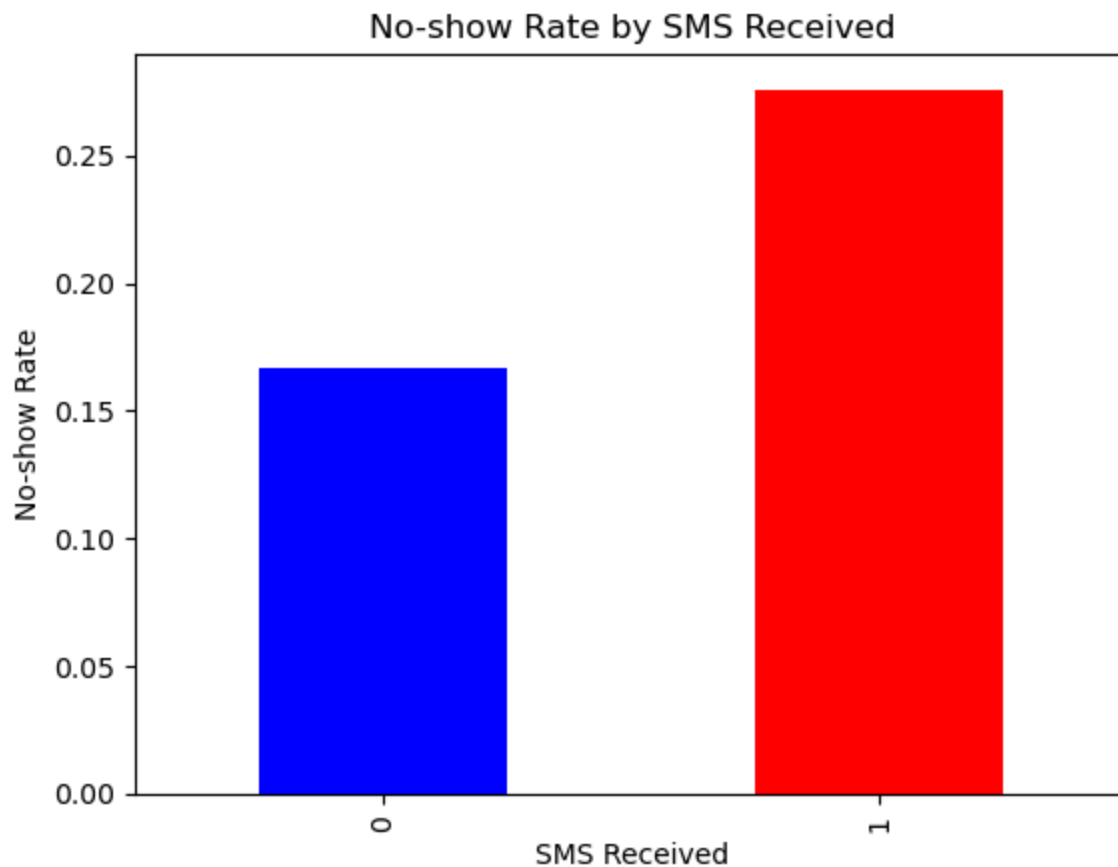
```
In [16]: # Calculating no-show rate by scholarship
no_show_rate_scholarship = df.groupby('Scholarship')['No-show'].apply(lambda x: (x

# Plotting the no-show rate by scholarship
no_show_rate_scholarship.plot(kind='bar', color=['orange', 'green'])
plt.title('No-show Rate by Scholarship')
plt.xlabel('Scholarship')
plt.ylabel('No-show Rate')
plt.show()
```



```
In [17]: # Calculating no-show rate by SMS received
no_show_rate_sms = df.groupby('SMS_received')['No-show'].apply(lambda x: (x == 'Yes').mean())


# Plotting the no-show rate by SMS received
no_show_rate_sms.plot(kind='bar', color=['blue', 'red'])
plt.title('No-show Rate by SMS Received')
plt.xlabel('SMS Received')
plt.ylabel('No-show Rate')
plt.show()
```



```
In [18]: # Selecting numerical columns for correlation
numerical_columns = ['Age', 'Scholarship', 'Hipertension', 'Diabetes', 'Alcoholism']

# Calculating the correlation matrix
corr_matrix = df[numerical_columns].corr()

# Plotting the heatmap
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```



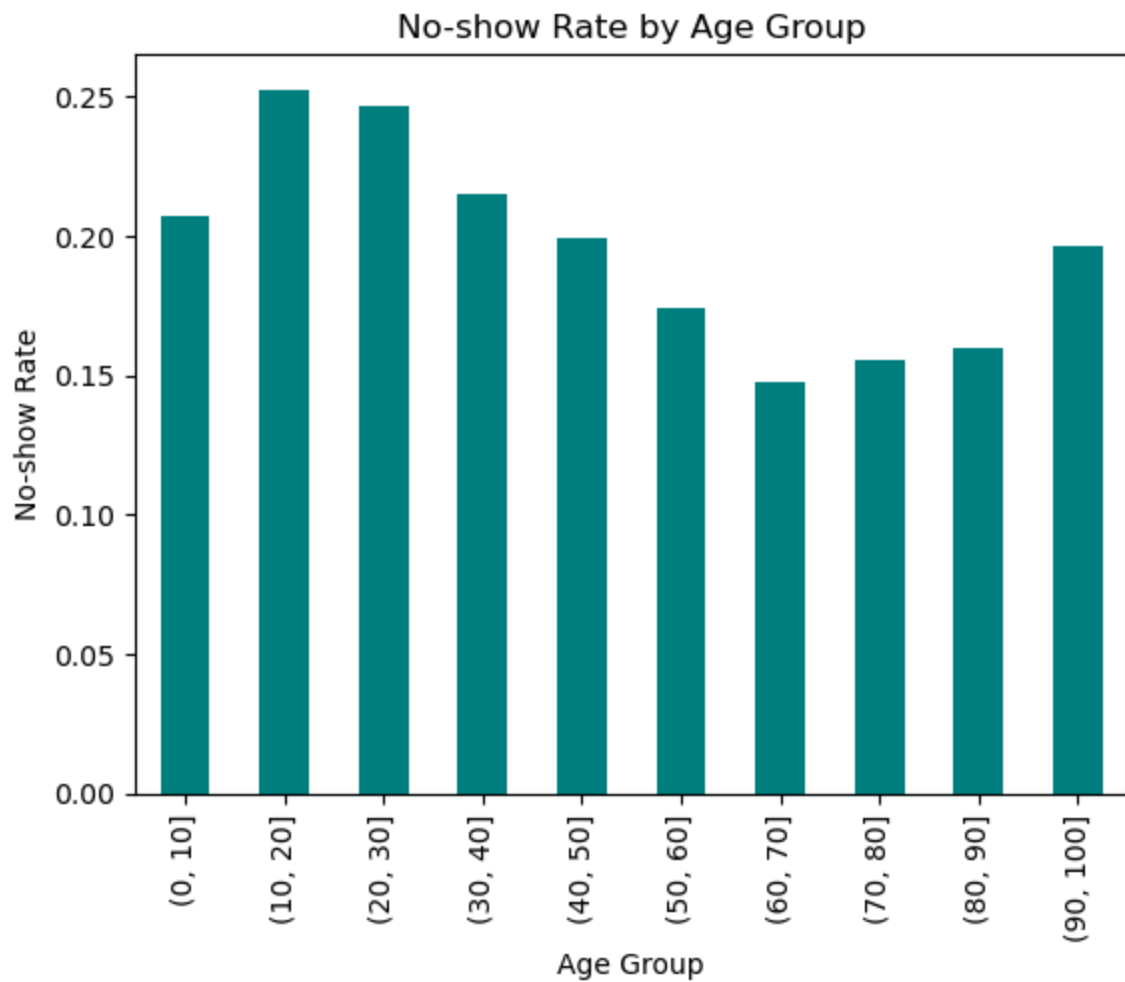
```
In [19]: # Creating age groups
df['AgeGroup'] = pd.cut(df['Age'], bins=range(0, 101, 10))

# Calculating no-show rate by age group
no_show_rate_age = df.groupby('AgeGroup')['No-show'].apply(lambda x: (x == 'Yes').mean())

# Plotting the no-show rate by age group
no_show_rate_age.plot(kind='bar', color='teal')
plt.title('No-show Rate by Age Group')
plt.xlabel('Age Group')
plt.ylabel('No-show Rate')
plt.show()
```

C:\Users\Administrator\AppData\Local\Temp\ipykernel\_11032\2774455716.py:5: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
no_show_rate_age = df.groupby('AgeGroup')['No-show'].apply(lambda x: (x == 'Yes').mean())
```



## Feature Engineering - Selection, Cleaning

### Cleaning

```
In [22]: # Check for missing values
print(df.isnull().sum())

# Remove duplicates
df = df.drop_duplicates()
```

```
PatientId          0
AppointmentID      0
Gender             0
ScheduledDay       0
AppointmentDay     0
Age               0
Neighbourhood      0
Scholarship        0
Hypertension       0
Diabetes           0
Alcoholism         0
Handcap           0
SMS_received       0
No-show            0
AgeGroup          3547
dtype: int64
```

```
In [23]: # Convert ScheduledDay and AppointmentDay to datetime
df['ScheduledDay'] = pd.to_datetime(df['ScheduledDay'])
df['AppointmentDay'] = pd.to_datetime(df['AppointmentDay'])

# Check unique values in categorical columns
print(df['Gender'].unique())
print(df['No-show'].unique())

['F' 'M']
['No' 'Yes']
```

## Create New Features

```
In [25]: # Calculate waiting time in days
df['WaitingTime'] = (df['AppointmentDay'] - df['ScheduledDay']).dt.days

# Extract day of the week from AppointmentDay
df['AppointmentDayOfWeek'] = df['AppointmentDay'].dt.day_name()

# Bin Age into groups
df['AgeGroup'] = pd.cut(df['Age'], bins=[0, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100])

# One-hot encode categorical variables
df = pd.get_dummies(df, columns=['Gender', 'Neighbourhood', 'AppointmentDayOfWeek',

# Encode No-show column
df['No-show'] = df['No-show'].apply(lambda x: 1 if x == 'Yes' else 0)

# Drop unnecessary columns
df = df.drop(['PatientId', 'AppointmentID', 'ScheduledDay', 'AppointmentDay'], axis
```

```
In [26]: df.shape
```

```
Out[26]: (110527, 104)
```

# Data Splitting

```
In [28]: from sklearn.model_selection import train_test_split

# Define features (X) and target (y)
X = df.drop('No-show', axis=1)
y = df['No-show']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
```

```
In [29]: from sklearn.preprocessing import StandardScaler

# Initialize the scaler
scaler = StandardScaler()

# Scale numerical features
X_train[['Age', 'WaitingTime']] = scaler.fit_transform(X_train[['Age', 'WaitingTime']])
X_test[['Age', 'WaitingTime']] = scaler.transform(X_test[['Age', 'WaitingTime']])
```

```
In [30]: X_train.shape
```

```
Out[30]: (88421, 103)
```

```
In [31]: y_train.shape
```

```
Out[31]: (88421,)
```

## Model selection - best fit analysis

```
In [33]: logistic = LogisticRegression(max_iter=100000)
logistic.fit(X,y)

print("Accuracy score on training {:.5f}".format(logistic.score(X_train,y_train)))
print("Accuracy score on testing {:.5f}".format(logistic.score(X_test,y_test)))
```

Accuracy score on training 0.79776

Accuracy score on testing 0.79929

```
In [34]: logistic100 = LogisticRegression(C=10000000, max_iter=100000)
logistic100.fit(X,y)

print("Accuracy score on training {:.5f}".format(logistic100.score(X_train,y_train)))
print("Accuracy score on testing {:.5f}".format(logistic100.score(X_test,y_test)))
```

Accuracy score on training 0.79776

Accuracy score on testing 0.79929

```
In [35]: logistic001 = LogisticRegression(C=0.001, max_iter=100000)
logistic001.fit(X,y)
```

```
print("Accuracy score on training {:.5f}".format(logistic001.score(X_train,y_train))
print("Accuracy score on testing {:.5f}".format(logistic001.score(X_test,y_test)))
```

Accuracy score on training 0.79776

Accuracy score on testing 0.79929

```
In [36]: svm = LinearSVC(max_iter=100000)
svm.fit(X,y)

print("Accuracy score on training {:.5f}".format(svm.score(X_train,y_train)))
print("Accuracy score on testing {:.5f}".format(svm.score(X_test,y_test)))
```

Accuracy score on training 0.79776

Accuracy score on testing 0.79929

```
In [37]: svm100 = LinearSVC(C=100, max_iter=100000)
svm100.fit(X,y)

print("Accuracy score on training {:.5f}".format(svm100.score(X_train,y_train)))
print("Accuracy score on testing {:.5f}".format(svm100.score(X_test,y_test)))
```

Accuracy score on training 0.79776

Accuracy score on testing 0.79929

```
In [38]: svm001 = LinearSVC(C=0.001, max_iter=100000)
svm001.fit(X,y)

print("Accuracy score on training {:.5f}".format(svm001.score(X_train,y_train)))
print("Accuracy score on testing {:.5f}".format(svm001.score(X_test,y_test)))
```

Accuracy score on training 0.79776

Accuracy score on testing 0.79929

```
In [39]: y_pred_logistic = logistic.predict(X_test)
y_pred_logistic
```

Out[39]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

```
In [40]: y_pred_svm = svm.predict(X_test)
y_pred_svm
```

Out[40]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

```
In [78]: import mglearn
from mglearn.plot_helpers import discrete_scatter
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
import numpy as np

# Assuming X and y are defined as NumPy arrays
# Select two features for visualization (e.g., columns 0 and 1)
X_plot = X[:, [0, 1]] # Use NumPy slicing to select two features

# Create a figure with two subplots
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
```



```

# Define the models
models = [LogisticRegression(max_iter=100000), LinearSVC(max_iter=100000)]

# Iterate over models and axes
for model, ax in zip(models, axes):
    # Fit the model using only the two selected features
    clf = model.fit(X_plot, y)

    # Plot the decision boundary using mglearn
    mglearn.plots.plot_2d_separator(clf, X_plot, fill=False, eps=0.5, ax=ax, alpha=

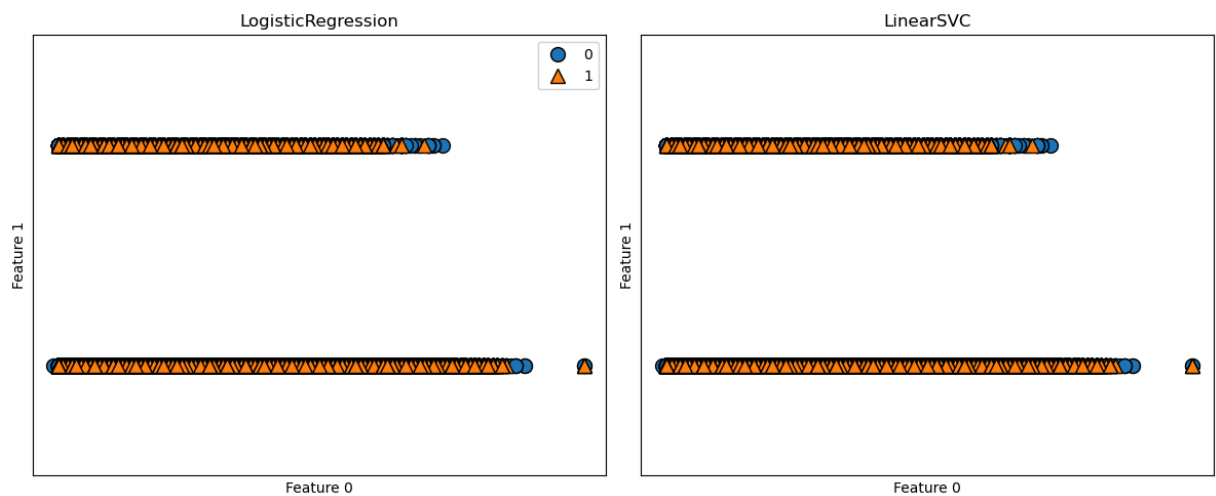
    # Scatter plot of the data points
    discrete_scatter(X_plot[:, 0], X_plot[:, 1], y, ax=ax)

    # Set title and labels
    ax.set_title("{}".format(clf.__class__.__name__))
    ax.set_xlabel("Feature 0")
    ax.set_ylabel("Feature 1")

# Add a Legend to the first subplot
axes[0].legend(loc='best')

# Adjust layout for better spacing
plt.tight_layout()
plt.show()

```



In [ ]: